

47% completed

Search Course

The Basics

Introduction

Program vs Process vs Thread

Concurrency vs Parallelism

Cooperative Multitasking vs Preemptive Multitasking

Synchronous vs Asynchronous

I/O Bound vs CPU Bound

Throughput vs Latency

Critical Sections & Race Conditions

Deadlocks, Liveness & Reentrant Locks

Mutex vs Semaphore

Mutex vs Monitor

Java's Monitor & Hoare vs Mesa Monitors

Semaphore vs Monitor

Amdahl's Law

Moore's Law

Multithreading in Java

Practice Mock Interview

Java Multithreading for Senior Engineering Interviews / ... / Program vs Process vs Thread

Program vs Process vs Thread

This lesson discusses the differences between a program, process and a thread. Also included is an example of a thread-unsafe program.

Program

A program is a set of instructions and associated data that resides on the disk and is loaded by the operating system to perform some task. An executable file or a python script file are examples of programs. In order to run a program, the operating system's kernel is first asked to create a new process, which is an environment in which a program executes.

Process

A process is a program in execution. A process is an execution environment that consists of instructions, user-data, and system-data segments, as well as lots of other resources such as CPU, memory, address-space, disk and network I/O acquired at runtime. A program can have several copies of it running at the same time but a process necessarily belongs to only one program.

Thread

Thread is the smallest unit of execution in a process. A thread simply executes instructions serially. A process can have multiple threads running as part of it. Usually, there would be some state associated with the process that is shared among all the threads and in turn each thread would have some state private to itself. The globally shared state amongst the threads of a process is visible and accessible to all the threads, and special attention needs to be paid when any thread tries to read or write to this global shared state. There are several constructs offered by various programming languages to guard and discipline the access to this global state, which we will go into further detail in upcoming lessons.

Caveats

Note a program or a process are often used interchangeably but most of the times the intent is to refer to a process.

There's also the concept of "multiprocessing" systems, where multiple processes get scheduled on more than one CPU. Usually, this requires hardware support where a single system comes with multiple cores or the execution takes place in a cluster of machines. Processes don't share any resources amongst themselves whereas threads of a process can share the resources allocated to that particular process, including memory address space. However, languages do provide facilities to enable inter-process communication.

Process

Global Variables

Thread

local Variables

Code

Thread

local Variables

Code

Thread

local Variables

Code

Counter Program

Below is an example highlighting how multi-threading necessitates caution when accessing shared data amongst threads. Incorrect synchronization between threads can lead to wildly varying program outputs depending on in which order threads get executed. Consider the below snippet of code

```
1. int counter = 0;
2.
3. void incrementCounter() {
4.     counter++;
5. }
```

The increment on **line 4** is likely to be decompiled into the following steps on a computer:

- Read the value of the variable counter from the register where it is stored
- Add one to the value just read

- Store the newly computed value back to the register

The innocuous looking statement on **line 4** is really a three step process! Now imagine if we have two threads trying to execute the same function `incrementCounter` then one of the ways the execution of the two threads can take place is as follows:

Lets call one thread as **T1** and the other as **T2**. Say the counter value is equal to 7.

1. **T1** is currently scheduled on the CPU and enters the function. It performs step A i.e. reads the value of the variable from the register, which is 7.
2. The operating system decides to context switch **T1** and bring in **T2**. **T2** gets scheduled and luckily gets to complete all the three steps **A**, **B** and **C** before getting switched out for **T1**. It reads the value 7, adds one to it and stores 8 back.
3. **T1** comes back and since its state was saved by the operating system, it still has the stale value of 7 that it read before being context switched. It doesn't know that behind its back the value of the variable has been updated. It unfortunately thinks the value is still 7, adds one to it and overwrites the existing 8 with its own computed 8. If the threads executed serially the final value would have been 9.

The problems should be apparent to the astute reader. Without properly guarding access to mutable variables or data-structures, threads can cause hard to find to bugs.

Since the execution of the threads can't be predicted and is entirely up to the operating system, we can't make any assumptions about the order in which threads get scheduled and executed.

Thread unsafe class

Take a minute to go through the following program. It increments a counter and decrements it an equal number of times. The final value of the counter should be zero, however, if you run the program enough times, you'll sometimes get the correct zero value, and at others, you'll get a non-zero value. We sleep the threads to give them a chance to run in a non-deterministic order.



```
1 import java.util.Random;
2
3 class DemoThreadUnsafe {
4
5     // We'll use this to randomly sleep our threads
6     static Random random = new Random(System.currentTimeMillis());
7
8     public static void main(String args[]) throws InterruptedException {
9
10        // create object of unsafe counter
11        ThreadUnsafeCounter badCounter = new ThreadUnsafeCounter();
12    }
```

Example of Thread Unsafe Code

[← Back lesson](#)

[Completed](#) [Next →](#)